

Final project: Learning a Value Function with a Neural Network, and Using it as Terminal Cost in an Optimal Control Problem

Alessandro Wegher
n.m. 249819

University of Trento
Department of Industrial Engineering
Via Sommarive 14, 38123 Povo (TN), Italy

1 Introduction

Solving OCPs over long time horizons can be computationally very expensive and introduce quite high latency in real-time control of a robotic system. To improve the performance of the optimal-trajectory computation, we can rely on a model that approximates the function relating the initial state of the MPC $x(t_k)$ to the cost we would obtain by going optimally to the final state at the end of the horizon $x(t_k + N \cdot T_s)$.

In order to address this challenge, we train a simple neural network as a function-approximation structure that takes the initial state of the MPC x_0 as input and approximates the value of the terminal cost function $V_{nn}(x_0, N)$, also called the cost-to-go function. The idea is that running the trained model is faster than evaluating the whole policy each time, because the former implies only a function evaluation, compared to the iterative algorithms involved in computing the optimal trajectory with CasADi at runtime.

This report presents the results of this approach for two non-linear mechanical systems: a single and a double pendulum. Each system has four main scripts, mirrored under `Final_Project/Single_Pend/` and `Final_Project/Double_Pend/`:

- **Training script** (`SinPendTrain.py` / `Double_pend_training.py`): generates the dataset by solving many OCPs in parallel from random initial-state configurations, then trains and saves a neural-network model for later use.
- **Test script** (`SinPendTest.py` / `Double_pend_Test.py`): used to test the performance of the trained model against ground-truth OCP costs.
- **Open-loop comparison script** (`SinPendComp.py` / `Double_pend_comp.py`, with multi-process variants): compares the *open-loop* cost of three MPC formulations over K random initial states (horizon M without terminal cost, horizon M with the NN terminal cost, horizon $M + N$ without terminal cost).
- **Closed-loop MPC script** (`SinPend_MPC.py` / `Double_pend_MPC.py`): actually drives the pendulum in receding-horizon fashion, for the same three formulations, from one fixed initial state, logging and plotting the resulting $q/\dot{q}/\tau$ trajectories. This is the closed-loop counterpart of the open-loop comparison above, built on a shared engine, `MPC_solve.py`.

An additional script, `OCP_solve.py`, defines the cost-computation procedure, which solves the OCP in Bolza's form for both the case with and without a terminal cost term (also referred

to as Mayer’s term, $\Phi(x_N)$):

$$\begin{aligned} & \text{minimize} && \sum_{i=0}^{N-1} (w_p \|q_i\|^2 + w_v \|\dot{q}_i\|^2 + w_a \|u_i\|^2) + \Phi(x_N) \\ & \text{subject to} && x \in X, u \in U \\ & && x_0 = x_{\text{init}} \\ & && x_{i+1} = x_i + f(x_i, u_i) \cdot \Delta t, \quad i = 0 \dots N-1 \end{aligned}$$

The dynamic constraint is numerically implemented with a simple explicit Euler method (first order); other schemes (Crank–Nicolson, RK4, Dormand–Prince) could be used instead, at the cost of extra implementation complexity, e.g., for stiffer dynamics.

The file `double_pend_dynamics.mw` is a Maple worksheet where the equations of motion of the double pendulum in state-space form were derived symbolically with the Euler–Lagrange approach. `neural_network.py` contains the `NeuralNetwork` class used to instantiate the regressor in the training scripts.

2 The Three MPC Formulations

The three formulations compared throughout this report are:

1. **Case 1 – M , no terminal cost:** a short prediction horizon M , myopic MPC.
2. **Case 2 – $M + N$, no terminal cost:** a long prediction horizon $M + N$, treated as the near-optimal reference (a longer horizon better approximates the infinite-horizon, globally optimal policy).
3. **Case 3 – $M + \text{NN terminal cost}$:** the same short horizon M as case 1, plus the learned terminal cost $V_{nn}(x_0, N) \cdot J_{\text{MAX}}$ appended to the Bolza cost, where J_{MAX} rescales the network’s normalized $[0, 1]$ output back to the true cost scale. The intent is for case 3 to approximate case 2’s trajectory quality at a fraction of case 2’s per-solve computational cost.

J_{MAX} is calibrated once per experiment as the maximum average cost-per-step, $J(x_0)/(M+N)$, observed over $K = 50$ random initial states solved with the long horizon $M+N$ (see `MPC_solve.compute_jmax`).

All three are legitimately called *MPC* formulations because they (i) solve a predictive OCP over a finite horizon, (ii) use the deterministic dynamics model to forecast future behaviour, and (iii) are meant to run recursively in a real-time controller. This report evaluates them in two complementary ways:

- **Open-loop, aggregate cost:** for each of $K = 50$ random initial states, solve each formulation *once* and compare the resulting optimal cost (Figs. 3 and 8).
- **Closed-loop, single trajectory** (new in this revision): fix one initial state and actually run each formulation in receding horizon for $T_{\text{sim}} = 5 \text{ s}$ ($\Delta t = 0.01 \text{ s}$, so 500 re-solves per case), re-solving at every step, applying only the first control, warm-starting the next solve, and logging the real $q(t)$, $\dot{q}(t)$, $\tau(t)$ trajectories (Figs. 4–10).

Matching the aggregate open-loop cost across many random states does not automatically imply that two formulations produce the same closed-loop trajectory for any single instance – the second evaluation mode is what actually certifies (or not) that the learned terminal cost behaves like a real terminal cost inside a running controller, not just in expectation.

3 Single Pendulum

The dynamics of this 1-DOF system are formulated using the Newton–Euler approach, implemented in the `SimPend` function, which returns a CasADi function used for numerical integration of the system dynamics within the OCP.

To generate a dataset of initial states $[q_0, \dot{q}_0]^T$, a 60×60 grid is built with `linspace`. Multithreading (Python’s `multiprocessing`) is used during dataset generation, enabling a larger dataset to be built in the same wall-clock time compared to a single-threaded approach. The dataset uses an OCP horizon of $N = 200$ time steps, solving 3600 independent OCPs.

3.1 Neural Network Structure

- Input layer: $x \in \mathbb{R}^2 \rightarrow [q_0, \dot{q}_0]^T$
- Hidden layers: 3 fully connected layers, each with 64 neurons
- Activation function: $\tanh(\cdot)$
- Output layer: $J(x) \in \mathbb{R}$, the predicted (normalized) terminal cost

3.2 Hyperparameters


- Learning rate: 0.008
- Loss function: Mean Squared Error
- Training dataset size: 3600 points (60×60 grid)
- Input range: $q_0 \in [-\pi, \pi]$ rad, $\dot{q}_0 \in [-\pi, \pi]$ rad/s
- Batch size: full batch (entire dataset, no held-out test batch)
- Output scaling: normalized cost over the maximum valid cost in the dataset, J_{MAX}
- Training epochs: 2000
- Lagrange running-cost weights: $[w_p, w_v, w_a]^T = [1, 0.1, 0.1]^T$

The trained model is saved and later used to generate colormap plots and validate its performance in `SinPendTest.py`. The final training loss is ≈ 0.002 .


3.3 Open-Loop MPC Cost Comparison

The comparison scripts `SinPendComp.py` and `SinPendComp.multiprocess.py` implement the three MPC formulations and compare their open-loop cost over $K = 50$ random initial states. We compute the difference between case 3 and case 2, normalized over the long-horizon (case 2) cost, to clearly evaluate the model’s performance – overall quite good, aside from a few outliers.


As expected, cases 2 and 3 yield similar cost values (Fig. 3, right), and these are consistently lower than the short-horizon case 1. This is coherent with intuition: a shorter horizon makes the solver greedier, since it has limited predictive capability, so control actions that look optimal in the short term can be suboptimal over a longer horizon. The lowest achievable cost would correspond to the infinite-horizon case, globally optimal and consistent across all MPC iterations.



(a) Cost distribution of the training dataset



(b) Loss over the training epochs



(c) Colormaps of the predicted (normalized) cost, computed cost, and normalized computed cost

Figure 1: Neural network training results for the single pendulum.




Figure 2: Normalized residuals plot for the single pendulum case.




Figure 3: Left: normalized cost difference between case 3 (NN terminal) and case 2 (long horizon). Right: average cost per step of the three MPC formulations across $K = 50$ random initial states.

Single pendulum closed-loop MPC, $x_0 = [0.30670429 \ 2.70414933]$




Figure 4: Closed-loop position q_1 and velocity $\dot{q}_1$ for the three MPC formulations, single pendulum.

Single pendulum closed-loop MPC, $x_0 = [0.30670429 \ 2.70414933]$




Figure 5: Closed-loop applied torque u_1 for the three MPC formulations, single pendulum.

3.4 Closed-Loop MPC Trajectories

To verify that the open-loop cost agreement above actually translates into comparable *control behaviour*, `SinPend_MPC.py` runs all three formulations in real receding-horizon closed loop, from the same fixed, seeded initial state $x_0 = [0.307, 2.704]^T$ (rad, rad/s), for $T_{\text{sim}} = 5$ s at $\Delta t = 0.01$ s (500 steps), warm-starting every re-solve. J_{MAX} was calibrated to 7.73 for this run.

The closed-loop result is more nuanced than the aggregate open-loop comparison suggests. Case 1 (short horizon, no terminal cost) and case 3 (short horizon, NN terminal cost) produce *almost indistinguishable* trajectories throughout the whole simulation, in both position/ velocity and torque – the two curves overlap essentially everywhere in Fig. 4–5. Case 2 (long horizon, no terminal cost) is instead the visible outlier: it commands a markedly larger initial torque (≈ -3.2 Nm vs. ≈ -2.4 Nm for cases 1/3), reaches a larger negative velocity dip, and drives q_1 to zero noticeably faster.

This is a meaningful counterpoint to Sec. 3.4: matching the *aggregate* open-loop cost across $K = 50$ random states (Fig. 3) does not guarantee that the learned terminal cost reproduces the long-horizon formulation’s actual closed-loop *policy* for a specific initial condition. Here, the terminal-cost gradient added to the $M = 20$ -step Bolza cost is evidently not strong enough to pull the short-horizon solver away from its myopic solution and toward the more foresighted behaviour that an explicitly longer horizon (case 2) provides – the network correctly predicts the *magnitude* of the cost-to-go on average, but that alone does not shape the short-horizon optimizer’s early control decisions the way a genuinely longer horizon does. A larger terminal-cost weight w_{term} , or a network trained to also match the *gradient* of the true value function (not just its value), are natural next steps to close this gap. (In terms of raw compute, case 3 still closed the loop in 8.1 s of wall-clock time here versus 11.5 s for case 2 – only a modest $1.4\times$ speed-up compared to the double pendulum below, since the single-pendulum horizons are already small enough that the solve-time gap between M and $M+N$ is much less pronounced.)

4 Double Pendulum

We consider a 2-DOF, fully actuated, nonlinear mechanical system. As in the single-pendulum case, the dynamics are integrated with explicit Euler; here, however, the equations of motion are derived via the Euler–Lagrange formulation, since the actuators act directly on the generalized coordinates, giving a straightforward computation of the generalized forces.

The full derivation is in the Maple file `double_pend.dynamics.mw`, where the state-space

equations are also computed. To validate the dynamics implementation, an example (natural-response, zero-input) trajectory starting from a random initial state is shown in Fig. 6.




Figure 6: Example zero-input (natural response) trajectory, validating the double-pendulum dynamics implementation.

To generate the initial conditions and compute the associated cost values for training the neural network, a custom `Dataset` class handles both the state generation and the parallelized cost computation. The `Pendulum` class parses the double-pendulum model parameters (link lengths, rotational inertias about the actuated axes, distances from the rotation axes to the centers of mass, and masses) from its URDF, including the Huygens–Steiner contributions to the inertia terms. A new, deeper architecture, `ExtendedNeuralNetwork`, extends the capacity of the single-pendulum `NeuralNetwork` class. The training dataset has $N_{\text{OCPs}} = 3000$ points.

4.1 Neural Network Structure


- Input layer: $x \in \mathbb{R}^4$, the initial state $x = [q_1(0), q_2(0), \dot{q}_1(0), \dot{q}_2(0)]^T$
- Hidden layers: 6 fully connected layers, each with 128 neurons
- Activation function: $\tanh$
- Output layer: $J(x) \in \mathbb{R}$, the predicted terminal cost

4.2 Hyperparameters


- Learning rate: 0.009
- Loss function: Mean Squared Error (MSE)
- Training dataset size: 3000 points
- Input range: $\vec{q} \in D = [-\pi, \pi]^2$, $\dot{\vec{q}} \in F = [-2\pi, 2\pi]^2$
- Batch size: full batch
- Output scaling: normalized by the maximum cost in the dataset, J_{MAX}

- Training epochs: 1500
- Lagrange running-cost weights: $[w_p, w_v, w_a]^T = [1, 0.05, 0.03]^T$


The trained model is saved and later used to generate the scatter plot and validate performance in `Double_pend_Test.py`. The final loss (≈ 0.04) is higher than the single- pendulum case, consistent with the harder 4-D learning problem.




(a) Cost distribution of the training dataset



(b) Loss over the training epochs



(c) Scatter plot: predicted vs. computed cost



(d) Residuals plot

Figure 7: Neural network training results for the double pendulum.

4.3 Open-Loop MPC Cost Comparison

`Double_pend_comp.py` and `Double_pend_comp_multi.py` implement the same three-case comparison as the single-pendulum case.

4.4 Closed-Loop MPC Trajectories

`Double_pend_MPC.py` runs the same closed-loop experiment as Sec. 3.4, for $T_{\text{sim}} = 5$ s at $\Delta t = 0.01$ s, from the same seeded initial-state distribution, with horizons $M = 70$, $N = 200$ ($M + N = 270$).

Here the picture is markedly different from the single-pendulum case: all three formulations produce trajectories that are visually indistinguishable across the full 5 s window, in position, velocity, *and* torque, for both joints (Figs. 9, 10). Unlike Sec. 3.4, case 1 (myopic, $M = 70$ steps = 0.7 s lookahead) does not visibly diverge from cases 2/3 here. A likely explanation is that, for




Figure 8: Left: normalized cost difference between case 3 (NN terminal) and case 2 (long horizon). Right: average cost per step of the three MPC formulations, double pendulum.




Figure 9: Closed-loop positions q_1, q_2 and velocities $\dot{q}_1, \dot{q}_2$ for the three MPC formulations, double pendulum.




Figure 10: Closed-loop applied torques u_1, u_2 for the three MPC formulations, double pendulum.

this particular seeded initial state (a damped settling toward the origin, with no swing-up or actuation-limit maneuvering involved), a 0.7 s lookahead is already long enough that the greedy and far-sighted policies agree – the double pendulum’s M is proportionally similar to the single pendulum’s ($M/(M+N) \approx 26\%$ vs. 20%), but in absolute time it is $3.5\times$ longer, which matters more for how much of the settling transient the controller can already “see”.

When the three formulations do agree this closely, the value of the terminal cost is not visible in *trajectory shape* but in *computational cost*: in this run, case 2 ($N_{\text{pred}} = 270$) took 178.9 s of wall-clock time for the closed-loop simulation, while case 3 ($N_{\text{pred}} = 70$, with the learned terminal cost, $J_{\text{MAX}} = 0.599$) took only 51.7 s – a $3.46\times$ speed-up for control that is, for this initial condition, essentially identical. This is the cleanest demonstration in this report of the method’s actual payoff: the same closed-loop behaviour at a fraction of the per-step solve cost. Whether this speed-up comes with a trajectory-quality cost for *other* initial states – particularly ones nearer the training domain’s high-cost boundary, where Fig. 7d shows the network is least accurate – is a natural extension of this experiment (evaluating several seeded x_0 rather than one, as already done for the open-loop comparison in Fig. 8).

5 Discussion and Final Remarks

As already observed for the single pendulum, the (extended) neural-network architecture predicts the terminal cost function accurately within the training domain in both systems, enabling effective integration in an MPC framework – Figs. 3 and 8 confirm this in aggregate, open-loop terms. The double-pendulum learning task is harder, owing to the higher-dimensional input and stronger nonlinearity of both dynamics and cost function; this shows up in the residuals plot (Fig. 7d), where larger prediction errors appear in specific regions of the input space, likely those with steep cost gradients or near the domain boundaries – consistent with the higher final training loss (≈ 0.04 vs. ≈ 0.002 for the single pendulum).

The closed-loop experiments added in this revision (Secs. 3.4, 4.4) are the central addition of this report: they test the learned terminal cost inside an actual running controller, not just as a one-shot cost predictor. The central payoff the network is meant to deliver – case 3 reproducing case 2’s trajectory quality at case 1’s per-step computational cost – is not automatic just because case 3 matches case 2 in *aggregate open-loop cost*; the single-pendulum result in Sec. 3.4 is a concrete instance where it does not hold for a specific initial condition, which is exactly the kind of failure mode that an open-loop-only evaluation cannot reveal.

Possible improvements. From the colormap plot (Fig. 1c) and the double-pendulum residuals (Fig. 7d), the network struggles most near the boundaries of the training domain, precisely where the cost is highest and its gradient is steepest – a harder regime for a fixed sampling density, and one where a saturating activation like $\tanh$ compresses large output values, since $\tanh : \mathbb{R} \rightarrow (-1, 1)$. Switching to a non-saturating (or less saturating) activation such as ReLU, Leaky ReLU, Swish, or ELU could help the network better approximate large cost values in high-gradient regions. Non-uniform sampling (denser near the boundaries) and normalizing cost values before training could also help. More expressive architectures (residual or attention-based networks) could better capture sharp transitions in the cost landscape. Finally, the closed-loop analysis above suggests that *matching the value function alone is not sufficient* for the terminal cost to reproduce the long-horizon closed-loop policy: increasing w_{term} , or training the network against the value function’s *gradient* in addition to its value (e.g., via a Sobolev-training loss), are promising directions to make the terminal cost actively steer the short-horizon solver, not just match its expected cost.